

Building Scalable Video Understanding Benchmarks through Sports

Anonymous submission

Abstract

Existing benchmarks for evaluating long video understanding falls short on multiple aspects, either lacking in scale or quality of annotations. These limitations arise from the difficulty in collecting dense annotations for long videos (*e.g.* actions, dialogues, etc.), which are often obtained by manually labeling many frames per second. In this work, we introduce an automated **Annotation and Video Stream Alignment Pipeline** (abbreviated **ASAP**). We demonstrate the generality of ASAP by aligning unlabeled videos of four different sports (Cricket, Football, Basketball, and American Football) with their corresponding dense annotations (*i.e.* commentary) freely available on the web. Our human studies indicate that ASAP can align videos and annotations with high fidelity, precision, and speed. We then leverage ASAP’s scalability to create **LCric**, a large-scale long video understanding benchmark, with over 1000 hours of densely annotated long Cricket videos (with an average sample length of ~ 50 mins) collected at virtually **zero annotation cost**. We benchmark and analyze state-of-the-art video understanding models on **LCric** through a large set of compositional multi-choice and regression queries. Our strong human baseline indicates significant room for new research to explore, and we will publicly release code for ASAP and **LCric** to enable building the next generation of long video understanding datasets and models.

1 Introduction

Humans learn and master skills (*e.g.* playing guitar) by associating and reasoning over episodic memories captured over days, months and years of failed and successful attempts. Thus, building systems capable of understanding and reasoning over very long streams of visual data has been a long-standing and crucial problem in Computer Vision.

Progress towards **Long-horizon Video Understanding** (we refer as **LVU**) has been greatly limited by the lack of densely annotated data. Creating an LVU benchmark requires manually annotating videos frame-by-frame, which is incredibly tedious and hard to scale. This constraint has limited the length of existing densely-annotated video understanding benchmarks (Table 1) from a few seconds (Jang et al. 2017; Sigurdsson et al. 2016; Gupta et al. 2021; Xu et al. 2016) to a few minutes (Zeng et al. 2016; Krishna et al. 2017; Wu and Krahenbuhl 2021; Zhou, Xu, and Corso 2018; Gella, Lewis, and Rohrbach 2018; Bain et al. 2020).

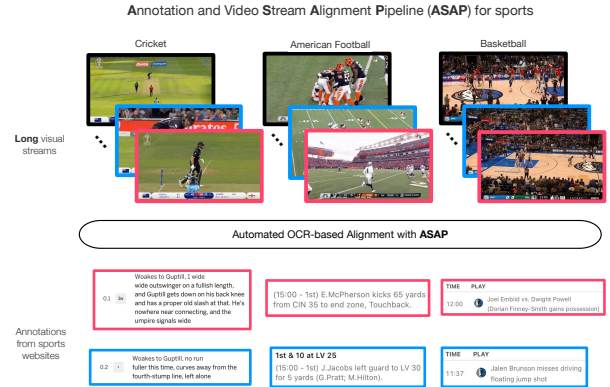


Figure 1: We propose the Annotation and Video Stream Alignment Pipeline (ASAP) for sports that uses an OCR based module to align unlabelled sports matches with structured annotations publicly available on the web. These newly aligned annotations can help compose structured queries that test for long-horizon video understanding skills.

A line of previous works (Huang et al. 2020; Tapaswi et al. 2016; Lei et al. 2018) in LVU have used readily available subtitles of TV shows or entire movies as dense annotations. While these videos are sufficiently long, manual annotations are still required to build non-trivial queries to evaluate LVU skills (Tapaswi et al. 2016; Lei et al. 2018), which greatly limits their scale. Another recent work (Wu and Krahenbuhl 2021) addresses this problem by extracting supervision from easily accessible YouTube metadata of nearly ~ 30000 movie clips spanning 1 – 3 minutes. However, the annotated clips are relatively short, and the proposed prediction tasks rely on noisy (and obscure) attributes (*e.g.* YouTube views, like-to-dislike ratio, and so on).

Sports matches are a rich source of long videos (*e.g.* a one-day Cricket match lasts nearly 8 hours) and usually have a brief scorecard embedded in the screen (as seen in Figure 1) that tracks the state of the match. Most sports matches also have dense annotations from experts available online (sports commentary describing major events in the game, *e.g.* (ESPN 2022b,a)). However, just the annotations or the videos are not helpful individually unless they are aligned with each other.

Therefore, we introduce ASAP, an automated annotation and video stream alignment pipeline, to automatically generate video datasets with dense annotations (i.e natural language commentary, major events in the match) by aligning arbitrarily long sports matches with their commentary freely available on the web (ESPN 2022b,a)). ASAP automatically parses metadata from the scorecard embedded in sports match videos using an OCR detector (Google 2022) and then uses the metadata to automatically align the videos with dense annotations available on the web. To demonstrate the generality of ASAP, we align unlabelled videos of four distinct sports (Cricket, Football/Soccer, Basketball, and American Football) with their corresponding web annotations, with an average of 95% of the annotations being aligned within ± 1 second of their occurrence in the video.

We then leverage ASAP’s scalability to create **LCric**, a large-scale LVU benchmark with 1008 hours of densely annotated Cricket videos at virtually zero annotation cost, by auto-labelling 131 cricket matches of average length 7.5 hours, containing nearly 475 timestamp recordings (balls per match) on average. To our knowledge, **LCric** is the first automatically-labelled sports video dataset that contains play-by-play annotations that span entire matches. To comprehensively evaluate LVU on **LCric**, we automatically curate multiple-choice (binary and n-way) and regression queries through simple composition with boolean operations. These queries are complex and require context aggregation ranging anywhere from 5 minutes to an hour of continuous playtime (video). In the past, such compositional query building has been leveraged in popular vision and language datasets (*e.g.* CLEVR (Johnson et al. 2017), GQA (Hudson and Manning 2019)).

We benchmark two recent state-of-the-art LVU models TQN (Zhang, Gupta, and Zisserman 2021a), MemViT (Wu et al. 2022a) on **LCric**, and find that their performance is significantly worse than our human baseline ($\sim 38\%$ drop on query reasoning accuracy when evaluated on very long clips containing ~ 50 minutes of playtime). This demonstrates significant room for new research to explore.

In summary, we make the following contributions – **1)** We propose ASAP, a fully automated and scalable video labeling pipeline for aligning videos of sports matches of four different sports (Cricket, Football, Basketball, and American Football) with dense annotations on the web. **2)** Using ASAP, we create **LCric**, a large-scale LVU benchmark with 1008 hours of densely annotated Cricket videos with virtually zero annotation cost. The size of **LCric** is only limited by our access to videos of Cricket matches, and we foresee the dataset being much larger given access to more videos. **3)** Finally, we benchmark the performance of two recent video understanding models on our dataset, provide ablations, and establish a human baseline on **LCric** to demonstrate significant room for modelling improvements on this benchmark. We will publicly release code for ASAP and **LCric** to enable future research on LVU.

Dataset	Avg. secs	clip	# Annotations	# Hours	Autolabel
VidSitu (2021)	10s		145K	81	\times
VideoStory (2018)	18s		123K	396	\checkmark
MSR-VTT (2016)	20s		200K	41	\times
Charades (2016)	30s		28K	82	\times
TGIF (2017)	30s		126K	86	\checkmark
TVQA (2018; 2019)	75s		<u>152K</u>	460	\times
VTW (2016)	90s		45K	213	\checkmark
MovieClips (2020)	120s		30K	1270	\checkmark
LVU (2021)	120s		11K	1270	\checkmark
YouCook II (2018)	<u>316s</u>		15K	176	\times
ActNet Captions (2017)	180s		100K	849	\times
LCric (ours)	2778s		62K	<u>1008</u>	\checkmark

Table 1: Comparison among annotated datasets for benchmarking video description and video understanding methods. **LCric** has an average clip length of ~ 2800 seconds, which is almost ten times larger than previous work, and has over 1000 hours of annotated videos.

2 Related Works

2.1 Existing benchmarks for LVU

(Wu and Krahenbuhl 2021) introduce the large-scale LVU benchmark built on movie clips and metadata publicly available on YouTube. However, the videos only range from 1-3 minutes, and the annotations are limited due to their dependence on YouTube metadata. (Oh et al. 2011) collect 29 hours of surveillance footage and bounding box annotations of major events but only have clips of length up to 3 minutes. (Corona et al. 2021) similarly collect 144 hours of surveillance footage by hiring actors to enact predefined scripts but only have clips of length up to 5 minutes. (Li et al. 2020) collect 430 videos, each 15 minutes long, and collect dense bounding box information for 80 different atomic actions. Though their videos are relatively long, our annotated videos are up to 45 minutes longer and are generated with no additional cost. (Cheng-Yang Fu and Berg 2017) collect the LoL dataset comprising 230 clips from the League of Legends video game, with each clip ranging from 30 to 50 minutes. However, they collect video highlight annotations based on very noisy and unreliable audience chat statistics. Video games also tend to have easy visual cues before major highlights that incentivize models to learn spurious correlations. Our tasks, by construction, force models to reason over a long horizon of events in a match.

2.2 Collecting dense annotations for videos

Annotating video datasets is extremely expensive. (Gupta et al. 2021; Xu et al. 2016; Sigurdsson et al. 2016) collect expensive annotations through Amazon Mechanical Turks (AMT) to label their clips with an associated text description, which greatly limits their scale (Table 1). Another line of work bootstraps from pre-existing annotations to generate new annotations. (Bain et al. 2020; Zeng et al. 2016; Gella, Lewis, and Rohrbach 2018) use pre-existing text descriptions and captions as annotations for their video clips. (Bain et al. 2020) use existing captions on YouTube and IMDb metadata to label 30000 movie clips, but assume these labels span the entirety of their clips. (Zeng et al. 2016) take user-generated titles as labels for 18100 user-generated clips, but

again assume that these labels span the entirety of their clips. (Gella, Lewis, and Rohrbach 2018) temporally align sentences from paragraph captions to social media videos to form annotated clips. Our dataset is densely annotated by temporally aligning publicly available sports annotations, which offer more structure than text descriptions and are therefore hierarchically composable, enabling the creation of queries that require large but dense context.

2.3 Video datasets based on sports

Recent interest in using computer vision to drive sports analytics (Tuyls et al. 2021) suggests the importance of a dense annotation pipeline for sports videos. Current methods for producing sports datasets involve some form of manual annotations. (Safdarnejad et al. 2015) manually label 4100 sports clips based on the given action. (Voeikov, Falaleev, and Baikulov 2020; Andriluka et al. 2017; Kazemi and Sullivan 2012) all use automatically-generated densely labelled pose annotations for sports videos but are not easily scalable because they run computationally expensive, frame-level models to generate their annotations. Larger datasets such as (Soomro, Zamir, and Shah 2012; Karpathy et al. 2014) exist but primarily focus on action recognition over a single clip, rather than a full sports video. Our dataset focuses on producing dense play-by-play annotations that span an entire sports match. Additionally, our general annotation pipeline can be easily extended to other sports, as most popular sports have rich commentary platforms on sports websites such as ESPN.

2.4 Video understanding models

Processing long videos is challenging, as it requires aggregating context over long horizons with limited computational and memory budgets. (Feichtenhofer et al. 2019) introduced SlowFast networks, which use a dual pathway operating at a low and high frame rate to enable the aggregation of context over longer horizons while capturing low-level visual attributes. (Feichtenhofer 2020) introduces a simple technique for progressive architecture expansion (along axes such as temporal, depth, width, etc.), inspired by feature selection in machine learning to achieve efficient models. Taking advantage of the implicit nature of transformers to handle long-range data (Bertasius, Wang, and Torresani 2021) proposes to adapt the standard transformer architecture for videos by enabling spatiotemporal feature learning directly from a sequence of frame-level patches. (Wu et al. 2022a) introduce MeMViT, a memory-augmented multi-scale vision Transformer, and greatly improves temporal support with minimal memory overhead and achieves state-of-the-art performance on a variety of video understanding benchmarks. While the trend shows the model’s capacity to handle longer and longer video clips more efficiently, the absence of a truly long-horizon dataset inhibits a fair comparison between these baselines and also inhibits the model’s transferability to real-world video understanding tasks.

2.5 Automated annotation pipelines

Automating annotation pipelines, even partially, is critical to developing large-scale datasets. (Ordonez, Kulkarni, and

Berg 2011) introduce the SBU dataset for image-text retrieval by pruning Flickr queries to pair queries with a set of images. (Sharma et al. 2018) introduce the Conceptual Captions dataset for image captioning by leveraging the “Alt-text” HTML attribute in web images. (Desai et al. 2021) introduce the RedCaps dataset with over 12 million image-text pairs harvested from curated sub-reddits. (Pont-Tuset et al. 2020) partially automate their annotation pipeline and collect multi-modal image annotations by asking annotators to describe an image through audio while simultaneously hovering their mouse over the region they are describing. We hope that our fully automated pipeline, ASAP, will help create long and densely annotated video datasets at an unprecedented scale.

3 ASAP: Annotation and Video Stream Alignment Pipeline for Sports Matches

Sports matches provide an abundant source of long videos, with matches often lasting multiple hours. These matches also have a rich source of corresponding play-by-play annotations (i.e. expert commentary, major events in the match) easily accessible on the web (ESPN 2022b,a). These play-by-play annotations are, however, not useful standalone as they are not aligned with the video of the match. To address this, we introduce ASAP, a fully automated annotation pipeline for automatically aligning sports videos with their corresponding play-by-play annotations available on the web. ASAP automatically parses metadata from the scorecard embedded in sports matches through an OCR detector and then uses the metadata to automatically align the videos with dense annotations available on the web. ASAP will enable us to create long video datasets with unprecedented scale with virtually zero annotation cost and is only limited by access to sports matches. We describe ASAP in more detail below.

3.1 Stage 1: Score Extraction and Eventwise Chunking

Scorecards contain metadata that help viewers understand where the match is at (*state*). For instance, Football, Basketball and American Football matches have precise information about play time, whereas, Cricket matches have information about the *ball* (Section 4.1) being played. For every frame in the match it is possible to crop out the desired scorecard (specified with a bounding box), and then extract the relevant metadata with the Google Cloud OCR API (shown in Figure 2).

However, we find that locating the scorecard in frames, as well as extracting metadata from it is non-trivial – due to the noisy and dynamic nature of the scorecard. We list some key challenges we encountered (and resolved) below.

Occlusion by advertisements, changing attributes of scorecards (format, shape, color) etc. To address issues posed by noisy and dynamic scorecards, we automatically generate a ‘reference frame’, a template image which is compared against all frames in the video with a simple L1 distance metric to remove spurious frames, i.e frames **1**) being occluded by advertisements or other miscellaneous

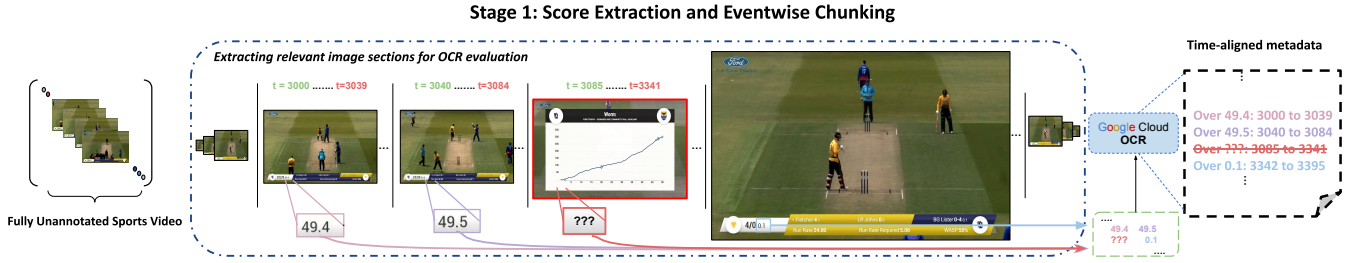


Figure 2: **Illustrating ASAP Stage 1 for Cricket matches.** ASAP automatically extracts the bounding box coordinates of the scorecard in a sports match and then extracts metadata from the scorecard by running OCR on each frame. Extracting metadata is often non-trivial. For instance, ASAP needs to automatically ignore one of the frames (marked in red) due to the scorecard being occluded by a miscellaneous patch showing players statistics.

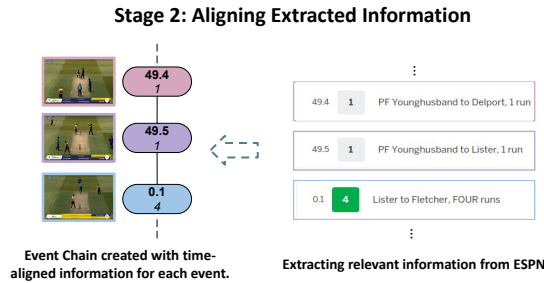


Figure 3: **Illustrating ASAP Stage 2 for Cricket matches.** ASAP scrapes dense play-by-play annotations from the web (<https://www.espnricinfo.com/>) and parses them into discrete events. ASAP then uses the extracted metadata from Stage 1 to align the annotations with the video of the match.

patches (shown in Figure 2), 2) having varying attributes of the scorecard (format, shape, color, etc.), and 3) having stale information on the scorecard. To automatically generate this ‘reference frame’, we first randomly sample frames from the middle of the match and use OCR calls to determine bounding boxes where text gradually changes over time (i.e the bounding box containing ‘0.1’ in Figure 2). We then use the corresponding bounding box generated by the OCR as the ‘reference frame’.

OCR call for every frame in a video can be expensive. Sports matches last for many hours and passing all frames of a 30 FPS video through the paid Google OCR API becomes prohibitively expensive. To reduce costs, we stack multiple scorecard boxes into a single image and get multiple annotations with a single OCR call. We also filter frames which contain the same metadata by detecting changes in the scorecard between consecutive frames using simple $L1$ distance.

3.2 Stage 2: Aligning dense annotations with videos

Dense play-by-play annotations (i.e expert commentary, major events in a match etc.) are often easily available on the web (ESPN 2022b,a). Since these play-by-play annotations are indexed by the metadata (precise *play time* or

ball information) extracted with the first stage of ASAP, we can align the videos with their corresponding annotations. In addition to aligning the annotations with the video, ASAP also processes the sequence of play-by-play annotations into a sequence of discrete events, which we refer to as an ‘event chain’ (Figure 3). While some sports (Cricket and Football) already contains discrete events (e.g. ‘foul’, ‘penalty’, ‘wicket’, ‘boundary’ etc.) in their annotation, for other sports (American Football and Basketball), we automatically parse the commentary and assign each play with a discrete event. These extracted event chains can then be used for easily evaluating LVU models. Models can be queried on different segments of the event chain to test for both short and long horizon reasoning. We concretely demonstrate the use of event chains for evaluation in Section 4.3.

How fast is ASAP? ASAP generates annotations with very high speed and requires just 10 minutes to align and process around 7 hours of video at 30FPS on a single machine. While the optimizations we introduce in Section 3.2 greatly improve annotation speed, methods such as multiprocessing, superior similarity matching metrics for filtering frames, and OCR predictions with a locally served OCR model could greatly increase annotation speed.

How accurate is ASAP? To verify ASAP’s ability to align dense annotations on the web with videos of sports matches, we conduct a study with human annotators on Amazon Mechanical Turk (AMT). We randomly sample clips from sports matches corresponding to 6 contiguous events in the event chain generated by ASAP. For all the generated clips, we ask human annotators to provide timestamps for all 6 events and then check whether the provided timestamps belong to the intervals generated for those events by ASAP. We plot the resulting accuracy of the timestamps in Figure 4 and find that ASAP is highly accurate, with an average accuracy of 95.3% across four very different sports, each differing in visual attributes, number of events and length of plays. The drop in accuracy for American football annotations can be explained by the inconsistency in the timestamps provided by ESPN. For regular plays, the timestamp indicates when the play begins; however, for ‘touchdowns’, the timestamp indicates when the team scored and not when the play begins. Additionally, penalties may affect the game clock, which we use to align our annotations, which some-




Figure 4: We find that human annotators rate annotations from ASAP to be highly accurate, with an average of 95% of the annotations being correctly aligned to the corresponding moment in a video (± 1 second) across four very different sports.

times leads to slight alignment issues for ASAP. We present more details in Appendix A.

4 Generating LCric dataset with ASAP

In this section, we describe how we leverage ASAP to build a long video understanding (LVU) dataset from Cricket videos online. We first provide a brief overview of the sport and describe the corresponding web annotations used by ASAP. Next, we describe a scalable approach for generating structured and compositional queries to evaluate LVU. Finally, we discuss important statistics of the queries and cricket videos in LCric.

4.1 Introduction to Cricket

We present a brief introduction to the sport of Cricket before introducing LCric. Cricket is played by two teams of 11 players each that alternate between *batting* and *fielding* throughout the game. The batting team aims to score **runs** by hitting a ball that is bowled from the fielding team out of the playing field. Meanwhile, the fielding team aims to prevent the batting team from scoring runs and dismiss all players in the batting team by taking their *wickets*. Each exchange where the fielding team bowls a valid ball and the batting team attempts to hit the ball to score runs is called a *ball* (or *delivery*) and a sequence of 6 *balls* is called an *over*. Each *ball* is an **atomic event** and there are 12 distinct possible events per ball, listed below:

- The batting team scores n runs, where $n \in \{0, \dots, 9\}$
- A wicket is taken and the current batsman is dismissed.
- A wide (invalid) ball is bowled, giving the batting team an extra run and another ball.

We present a detailed discussion of different phases of Cricket in Appendix B.

4.2 LCric: Overview

Leveraging ASAP’s scalability, we create LCric, a large-scale LVU benchmark with 1008 hours of densely annotated Cricket videos with virtually zero annotation cost, by

auto-labelling 131 cricket matches of average length ~ 7.5 hours, containing nearly 475 timestamp recordings (balls per match) on average. ASAP automatically labels all the balls in a Cricket match with 1 of 12 events (Section 4.1) to generate a sequence of events (i.e *event chain* shown in Figure 3) for a cricket match. We then generate annotated video clips by segmenting the videos along with the aligned event chain into a contiguous sequence of 10-over (~ 50 minutes) clips.

4.3 LCric: Evaluating LVU with compositional queries

Motivation: An LVU system needs to possess two types of skills: a) the ability to reliably detect local (short-term) events – *e.g.*, classifying an atomic event in Cricket (say wide, wicket, or run), and b) the ability to aggregate information across these local events given a task (which we refer to as a query) – *e.g.*, counting the total number of runs scored by the batting team in an arbitrarily long video. To test for these LVU skills using LCric, we automatically compose and filter binary and multi-choice queries, which are then paired with long video segments. For completeness, we also experiment with a regression query.

Min-Max occurrence query. This query helps to check occurrence of a single event, we first sample an atomic event from the set of all possible events, and then sample two numbers, o_{min} and o_{max} , to denote the minimum and maximum number of occurrences needed for this query to be *true*. An example of such a query is – for a given video, did a wide ball (an event) occur between 3 and 5 times inclusive?

Binary queries by chaining occurrence queries. To increase query diversity and complexity, we sample n_{chain} different min-max occurrence queries and combine them using [and]/[or] operators. For example, for a given video spanning 10-overs (~ 50 mins), “did a wide ball occur between 3 to 5 times [and] did a ball with 2 runs scored occur 1 to 3 times?”. All **binary** queries in LCric are formed by chaining 1-5 different min-max occurrence queries.

Multi-choice queries by counting occurrences. We expand upon the binary occurrence queries by generating multi-choice occurrence queries, which ask models to directly predict the number of occurrences rather than predicting membership in a range. An example of such a query is – for a given video, how many times a did a wide ball occur after a ball with 4 runs? Please note that these events are sequential, but not necessarily contiguous. As most non-trivial **multi-choice** events in LCric occur between 0-9 times in a given clip, we use $\{0, \dots, 9\}$ as our answer choices.

Filtering unbalanced queries. We can compose many LVU multi-choice and binary queries using the above formulation, however, not all queries are necessarily balanced. Due to the rarity of certain events occurring in Cricket, some queries are far easier to guess correctly than others. For example, in a 45 minutes clip (spanning 10 overs), the query – “did a ball with 2 runs occur between 0 to 10 times” is true with a probability of 87%. We filter such queries based on the probability of their occurrence in training matches, and ensure the average probability of occurrence of the selected queries to be between 0.45 – 0.55 to avoid bias.

Regression query for counting runs. Lastly, we also experiment with a single regression query that asks the model to predict the number of runs scored as a regression output for a given video sequence.

4.4 LCric: Statistics and Dataset Splits

Statistics The dataset currently includes 1008 hours of cricket match videos across 131 unique matches (average length of 7.5 hours), along with 61957 ball-by-ball annotations. All the videos are preprocessed with a resolution of 360p and we will provide links to the source videos, which are of a higher resolution.

Dataset splits To effectively test generalization, we split all the matches in **LCric** into train, validation and test splits and ensure a ‘60:20:20’ ratio of the number of hours in each split. Due to a limited computational budget, we present ablations on a subset of **LCric** and refer to this as **LCric-Mini**, which has around 420 hours of labelled Cricket matches, and therefore enables us to train ablations experiments in a shorter duration (2-3 days per experiment). We generate splits for **LCric-Mini** identically to **LCric**.

5 Experiments on LCric

In this section, we describe our experiments benchmarking state-of-the-art LVU models on **LCric**. We also provide a human baseline on **LCric** and demonstrate significant room for improvement. Finally, we analyze the performance of our baselines and present key insights to spur future modelling improvements on **LCric**.

5.1 Experimental Setup

Preprocessing: Following past work in LVU ((Zhang, Gupta, and Zisserman 2021b)), we process (sample) videos with a lower FPS to make training over long videos feasible. We process our longest clips (containing 10 overs of the match) at 0.1 FPS, and process clips of 2-8 overs at 0.5 FPS respectively. We **remove the scorecard** from all frames to prevent annotation leakage and process frames at a resolution of 128 x 128.

Evaluation Metrics: We compute the following evaluation metrics **1)** Classification accuracy for binary and multi-choice queries **2)** Average L1 norm for regression queries.

5.2 Baseline Models and Training Scheme

Previous works (Fan et al. 2021; Feichtenhofer et al. 2019; Feichtenhofer 2020) in LVU use pretrained CNNs (Bengio and Lecun 1997) and Transformers (Vaswani et al. 2017) paired with explicit memory modules for modelling long contexts. However, none of these methods can scale to video clips longer than a few minutes. Since our query set requires reasoning over contexts ranging up to an hour, we choose two recent state-of-the-art video understanding models, Temporal Query Network (TQN) (Zhang, Gupta, and Zisserman 2021b) and MeMViT (Wu et al. 2022b) as benchmarks – given their effective caching mechanisms for processing very long videos.

Temporal Query Network (TQN) (Zhang, Gupta, and Zisserman 2021b) TQN uses stochastic memory banks to




Figure 5: Accuracy (and avg. norm for regression) of various baselines on **LCric-Mini**, evaluated with our filtered set of binary, multi-choice, and regression occurrence queries for clips of over-length $\{2, 4, 6, 8, 10\}$. We find that our baseline LVU models severely degrade in performance with longer time horizons. Additionally, human performance far exceeds that of current models and therefore suggests significant room for modelling improvements on **LCric**.

efficiently model long horizon videos. They also introduce a transformer-based multi-query head to generate responses for multiple queries with a single pass through the network. TQN uses S3D (Xie et al. 2018) for the visual backbone network and processes non-overlapping contiguous sequence of 8 frames with a temporal stride of 1.

Memory-augmented Multiscale Vision Transformer (MeMViT) (Wu et al. 2022b) MeMViT applies a memory caching strategy by processing videos in an online fashion, allowing the model to efficiently store context to reason over a long horizon. MeMViT builds upon ViT (Dosovitskiy et al. 2020) by using a novel pooling method and a dynamic patch resolution approach to reduce computational costs while processing long clips. We adapt MeMViT to handle our multi-query setting by leveraging the multi-query head introduced in TQN.

Training Scheme We employ two different training schemes – **1) Homogeneous training** where we train different models for the three different types of queries (binary, multi-choice and regression) and **2) Mixed training** where we train a single model for all three types of queries.

5.3 Human Baseline

To quantify the room for modelling improvements on **LCric**, we benchmark the performance of human annotators through the AMT platform. We provide annotators with video clips from **LCric** and ask them to predict the sequence of ball-by-ball events (4.1). To compute human performance on our queries, we assume that given an event chain, humans can answer these queries by applying logical operators without mistakes.

5.4 Key Results

Performance of both TQN and MeMViT degrades rapidly and approaches random for very long clips To understand the impact of length of the videos on task performance, we train different baseline models for clips with over-lengths ranging from 2 overs (~ 10 minutes) to 10 overs

Model	Training scheme	BC Acc. \uparrow	MC Acc. \uparrow	R L1-Norm \downarrow
TQN	Mixed	57.68%	19.05%	17.21
MeMViT	Mixed	54.31%	16.71%	21.79
TQN	Homogeneous	60.74%	20.19%	10.63
MeMViT	Homogeneous	56.53 %	17.79%	<u>11.95</u>
Human Perf.	--	96.34 %	96.29%	0.215

Table 2: Baseline experiment results on the full **LCric** split with **10-over clips**. BC: Binary Choice Queries, MC: Multiple-Choice Query (10-choice), R L1-Norm: Avg. L1-Norm on Regression Queries. We find that TQN performs much better than MeMViT in different training schemes across different query types. We also note that training models under the Homogeneous training scheme greatly improves performance, especially for the regression query.

Model	Segmentation Scheme	BC Acc. \uparrow	MC Acc. \uparrow	R L1-Norm \downarrow
TQN	Ground Truth	80.34%	36.32%	6.89
TQN	Uniform	60.29%	15.18%	19.31
Human Perf.	--	96.34 %	96.29%	0.215

Table 3: To understand the impact of context aggregation on task performance, we first extract event chains from 10-over clips on **LCric-Mini** by running an event classifier on different segments of the clip. We then evaluate our queries against the generated event chains and find that even with access to ground truth event segments, performance is much worse than the human baseline. BC: Binary Choice Query Set (32 queries), MC: Multiple-Choice Query (10 queries), R L1-Norm: Avg. L1 Norm on Regression Queries.

(~ 50 minutes). Figure 5 shows that performance rapidly decreases with increasing clip length and approaches the random baseline for binary and multi-choice queries. This result, in addition to the strong human baseline, demonstrates significant room for modelling improvements.

Models need to effectively aggregate context to perform well on LCric To understand the importance of aggregating context for **LCric**, we first train a TQN ‘event classifier’ model to predict 1 of 12 events (Section 4.1) in a video clip. The model is trained on ground truth annotations from **LCric-Mini** and has a fairly high test accuracy of 84.79%. We then divide the clips into a contiguous sequence of ‘event segments’ by either using ground truth segmentations from **LCric-Mini** (labeled ‘Ground Truth’ in Table 3) or by uniformly dividing the clips into 60 contiguous segments (labeled ‘Uniform’ in Table 3), as each 10-over clip contains 60 events. Finally, we leverage the learned ‘event classifier’ model to generate event chains by sequentially predicting events on the ‘event segments’ and evaluate different queries on these event chains. We report performance in Table 3 and make two observations – **1)** Although access to ground truth segments of events leads to better performance, as it aids the event classifier in making more accurate predictions, the performance is still much worse than the human baseline. **2)** TQN model in Table 2, which aggregates context across 10-overs, performs much better than ‘Uniform’ in Table 3. Both these observations suggest that improving context aggregation is an important direction towards improving performance on **LCric**.

TQN performs better than MeMViT on all query types Table 2 shows that TQN performs much better than MeMViT across different query types and training schemes (+2% on binary and multi-choice queries). While we expected MeMViT to perform better due to the clever caching mechanism in its architecture, we think that adapting MeMViT with the TQN multi-query head could have caused the performance drop. We hope future work can better integrate the clever caching mechanism of MeMViT to the multi-query setting as it is infeasible to train one model for every new query.

Homogeneous training scheme leads to superior performance than Mixed training Our results in Table 2 indicate that training different models for different query types leads to much better performance than training a single model for all queries. We also observe that performance on the regression query improves by ~ 7 points and we posit that the model needs extra representational capacity to answer the regression queries with high precision.

6 Discussion

In this work, we introduce ASAP, a fully automated annotation and video stream alignment pipeline for sports matches. ASAP automatically aligns unlabelled videos of sports matches with corresponding dense annotations (*i.e.* commentary) freely available on the web. We demonstrate the generality of ASAP by aligning unlabelled matches of four very different sports (Cricket, Football/Soccer, Basketball, and American Football) with their corresponding annotations on the web. ASAP is highly accurate across the four sports (as judged by human annotators), and is therefore robust to varying visual attributes, number of events and length of plays. We then demonstrate ASAP’s potential to generate large-scale video datasets with **no additional annotation cost** by generating **LCric**, a large-scale long video understanding benchmark, with over 1000 hours of densely annotated long Cricket videos (with an average sample length of ~ 50 minutes). We extensively benchmark state-of-the-art LVU models and establish a human baseline on **LCric**. Our strong human baseline, coupled with poor performance of state-of-the-art models, validates **LCric** as an effective benchmark for the next generation of LVU models.

Our analysis suggests that effectively integrating clever memory/caching mechanisms in our multi-query setting could lead to much better performance on **LCric**. We also find that recent LVU models struggle to reason over very long contexts and we hope the next generation of architectures can accurately localize important events and effectively reason over a very long horizon. While we demonstrate ASAP’s potential by creating **LCric**, we hope future work extends, improves and leverages the various modules in ASAP to generate annotated video datasets for other sports and domains with an unprecedented scale and cost efficiency. We also hope future work can build upon our automated approach for evaluating LVU models on **LCric** and present alternate automated evaluation strategies for LVU.

References

- Andriluka, M.; Iqbal, U.; Milan, A.; Insafutdinov, E.; Pishchulin, L.; Gall, J.; and Schiele, B. 2017. PoseTrack: A Benchmark for Human Pose Estimation and Tracking. *CoRR*, abs/1710.10000.
- Bain, M.; Nagrani, A.; Brown, A.; and Zisserman, A. 2020. Condensed Movies: Story Based Retrieval with Contextual Embeddings. arXiv:2005.04208.
- Bengio, Y.; and Lecun, Y. 1997. Convolutional Networks for Images, Speech, and Time-Series.
- Bertasius, G.; Wang, H.; and Torresani, L. 2021. Is Space-Time Attention All You Need for Video Understanding? In *International Conference on Machine Learning*, 813–824. PMLR.
- Cheng-Yang Fu, M. B., Joon Lee; and Berg, A. C. 2017. Video Highlight Prediction Using Audience Chat Reactions. In *EMNLP*.
- Corona, K.; Osterdahl, K.; Collins, R.; and Hoogs, A. 2021. MEVA: A Large-Scale Multiview, Multimodal Video Dataset for Activity Detection. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, 1060–1068.
- Desai, K.; Kaul, G.; Aysola, Z.; and Johnson, J. 2021. RedCaps: Web-curated image-text data created by the people, for the people. In *NeurIPS Datasets and Benchmarks*.
- Dosovitskiy, A.; Beyer, L.; Kolesnikov, A.; Weissenborn, D.; Zhai, X.; Unterthiner, T.; Dehghani, M.; Minderer, M.; Heigold, G.; Gelly, S.; et al. 2020. An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale. *arXiv preprint arXiv:2010.11929*.
- ESPN. 2022a. ESPN Soccer Commentary. <https://www.espn.in/football/commentary>.
- ESPN. 2022b. ESPNCricinfo. www.espncriinfo.com/.
- Fan, H.; Xiong, B.; Mangalam, K.; Li, Y.; Yan, Z.; Malik, J.; and Feichtenhofer, C. 2021. Multiscale vision transformers. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 6824–6835.
- Feichtenhofer, C. 2020. X3D: Expanding Architectures for Efficient Video Recognition. In *Proc. CVPR*.
- Feichtenhofer, C.; Fan, H.; Malik, J.; and He, K. 2019. Slowfast networks for video recognition. In *Proceedings of the IEEE/CVF international conference on computer vision*, 6202–6211.
- Gella, S.; Lewis, M.; and Rohrbach, M. 2018. A Dataset for Telling the Stories of Social Media Videos. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, 968–974.
- Google. 2022. Google Cloud Optical Character Recognition. <https://cloud.google.com/vision/docs/ocr>.
- Gupta, A.; Yatskar, M.; Nevatia, R.; and Kembhavi, A. 2021. Visual Semantic Role Labeling for Video Understanding. In *CVPR 2021*.
- Huang, Q.; Xiong, Y.; Rao, A.; Wang, J.; and Lin, D. 2020. MovieNet: A Holistic Dataset for Movie Understanding. In *Proceedings of the European Conference on Computer Vision (ECCV)*.
- Hudson, D. A.; and Manning, C. D. 2019. Gqa: A new dataset for real-world visual reasoning and compositional question answering. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 6700–6709.
- Jang, Y.; Song, Y.; Yu, Y.; Kim, Y.; and Kim, G. 2017. Tgiffa: Toward spatio-temporal reasoning in visual question answering. In *Proc. CVPR*.
- Johnson, J.; Hariharan, B.; Van Der Maaten, L.; Fei-Fei, L.; Lawrence Zitnick, C.; and Girshick, R. 2017. Clevr: A diagnostic dataset for compositional language and elementary visual reasoning. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2901–2910.
- Karpathy, A.; Toderici, G.; Shetty, S.; Leung, T.; Sukthankar, R.; and Fei-Fei, L. 2014. Large-Scale Video Classification with Convolutional Neural Networks. In *2014 IEEE Conference on Computer Vision and Pattern Recognition*, 1725–1732.
- Kazemi, V.; and Sullivan, J. 2012. Using Richer Models for Articulated Pose Estimation of Footballers. In *BMVC*.
- Krishna, R.; Hata, K.; Ren, F.; Fei-Fei, L.; and Niebles, J. C. 2017. Dense-Captioning Events in Videos. In *International Conference on Computer Vision (ICCV)*.
- Lei, J.; Yu, L.; Bansal, M.; and Berg, T. L. 2018. TVQA: Localized, Compositional Video Question Answering. In *EMNLP*.
- Lei, J.; Yu, L.; Berg, T. L.; and Bansal, M. 2019. TVQA+: Spatio-Temporal Grounding for Video Question Answering. In *Tech Report, arXiv*.
- Li, A.; Thotakuri, M.; Ross, D. A.; Carreira, J.; Vostrikov, A.; and Zisserman, A. 2020. The ava-kinetics localized human actions video dataset. *arXiv preprint arXiv:2005.00214*.
- Oh, S.; Hoogs, A.; Perera, A.; Cuntoor, N.; Chen, C.-C.; Lee, J. T.; Mukherjee, S.; Aggarwal, J. K.; Lee, H.; Davis, L.; Swears, E.; Wang, X.; Ji, Q.; Reddy, K.; Shah, M.; Vondrick, C.; Pirsavash, H.; Ramanan, D.; Yuen, J.; Torralba, A.; Song, B.; Fong, A.; Roy-Chowdhury, A.; and Desai, M. 2011. A large-scale benchmark dataset for event recognition in surveillance video. In *CVPR 2011*, 3153–3160.
- Ordonez, V.; Kulkarni, G.; and Berg, T. 2011. Im2Text: Describing Images Using 1 Million Captioned Photographs. In Shawe-Taylor, J.; Zemel, R.; Bartlett, P.; Pereira, F.; and Weinberger, K. Q., eds., *Advances in Neural Information Processing Systems*, volume 24. Curran Associates, Inc.
- Pont-Tuset, J.; Uijlings, J.; Changpinyo, S.; Soricut, R.; and Ferrari, V. 2020. Connecting Vision and Language with Localized Narratives. In *ECCV*.
- Safdarnejad, S. M.; Liu, X.; Udpa, L.; Andrus, B.; Wood, J.; and Craven, D. 2015. Sports Videos in the Wild (SVW): A Video Dataset for Sports Analysis. In *Proc. International Conference on Automatic Face and Gesture Recognition*. Ljubljana, Slovenia.
- Sharma, P.; Ding, N.; Goodman, S.; and Soricut, R. 2018. Conceptual Captions: A Cleaned, Hypernymed, Image Alt-text Dataset For Automatic Image Captioning. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 2556–2565. Melbourne, Australia: Association for Computational Linguistics.
- Sigurdsson, G. A.; Varol, G.; Wang, X.; Farhadi, A.; Laptev, I.; and Gupta, A. 2016. Hollywood in homes: Crowdsourcing data collection for activity understanding. In *Proc. ECCV*.
- Soomro, K.; Zamir, A. R.; and Shah, M. 2012. UCF101: A Dataset of 101 Human Actions Classes From Videos in The Wild. *CoRR*, abs/1212.0402.

Tapaswi, M.; Zhu, Y.; Stiefelbogen, R.; Torralba, A.; Urtasun, R.; and Fidler, S. 2016. MovieQA: Understanding Stories in Movies through Question-Answering. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*.

Tuyts, K.; Omidshafiei, S.; Muller, P.; Wang, Z.; Connor, J.; Hennes, D.; Graham, I.; Spearman, W.; Waskett, T.; Steel, D.; Luc, P.; Recasens, A.; Galashov, A.; Thornton, G.; Elie, R.; Sprechmann, P.; Moreno, P.; Cao, K.; Garnelo, M.; Dutta, P.; Valko, M.; Heess, N.; Bridgland, A.; Pérolat, J.; De Vylder, B.; Eslami, S. M. A.; Rowland, M.; Jaegle, A.; Munos, R.; Back, T.; Ahamed, R.; Bouton, S.; Beau-guerlange, N.; Broshear, J.; Graepel, T.; and Hassabis, D. 2021. Game Plan: What AI can do for Football, and What Football can do for AI.

Vaswani, A.; Shazeer, N.; Parmar, N.; Uszkoreit, J.; Jones, L.; Gomez, A. N.; Kaiser, Ł.; and Polosukhin, I. 2017. Attention is all you need. In *Proc. NeurIPS*.

Voeikov, R.; Falaleev, N.; and Baikulov, R. 2020. TTNNet: Real-time temporal and spatial video analysis of table tennis. *CoRR*, abs/2004.09927.

Wu, C.-Y.; and Krahenbuhl, P. 2021. Towards long-form video understanding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 1884–1894.

Wu, C.-Y.; Li, Y.; Mangalam, K.; Fan, H.; Xiong, B.; Malik, J.; and Feichtenhofer, C. 2022a. MeMViT: Memory-Augmented Multiscale Vision Transformer for Efficient Long-Term Video Recognition. *arXiv preprint arXiv:2201.08383*.

Wu, C.-Y.; Li, Y.; Mangalam, K.; Fan, H.; Xiong, B.; Malik, J.; and Feichtenhofer, C. 2022b. MeMViT: Memory-Augmented Multiscale Vision Transformer for Efficient Long-Term Video Recognition. *arXiv preprint arXiv:2201.08383*.

Xie, S.; Sun, C.; Huang, J.; Tu, Z.; and Murphy, K. 2018. Rethinking spatiotemporal feature learning: Speed-accuracy trade-offs in video classification. In *Proc. ECCV*.

Xu, J.; Mei, T.; Yao, T.; and Rui, Y. 2016. MSR-VTT: A Large Video Description Dataset for Bridging Video and Language. In *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 5288–5296.

Zeng, K.-H.; Chen, T.-H.; Niebles, J. C.; and Sun, M. 2016. Title Generation for User Generated Videos. volume 9906. ISBN 978-3-319-46474-9.

Zhang, C.; Gupta, A.; and Zisserman, A. 2021a. Temporal Query Networks for Fine-grained Video Understanding. *arXiv preprint*.

Zhang, C.; Gupta, A.; and Zisserman, A. 2021b. Temporal query networks for fine-grained video understanding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 4486–4496.

Zhou, L.; Xu, C.; and Corso, J. J. 2018. Towards Automatic Learning of Procedures From Web Instructional Videos. In *AAAI Conference on Artificial Intelligence*.

A ASAP

A.1 Alignment Verification

We verify the accuracy of annotations made by ASAP by providing human annotators with a clip containing a contiguous sequence of events, and asking them to provide the timestamps in the video for when each event occurred. Additionally, all scorecard information is masked in each provided clip.

Verification of different sports For cricket, we built an AMT interface and asked annotators to provide both the timestamps and events that occurred in a clip for over 1200 events to verify that both the ASAP alignment process and video quality were sufficient, which we discuss more in Appendix B. For verifying and demonstrating the generality of ASAP pipeline, we annotate three different sports, namely, American football, football, and basketball, and verify it using a similar interface. Due to the limited mturk budget, we used two of the in-house annotators for the verification of these three sport’s annotations by providing the humans with clips from 6 hours of match footage for each sport and had them verify (by annotating) 240 events for each sport.

American Football Alignment Issues We note that the reason why the verification accuracy for American football in Figure 4 is lower than the other sports is because for most standard plays, the timestamps provided are for when the play started. However, when a team scores or is given a penalty, the timestamp provided for the next play is either the end of the play, or when it happened. We were only able to have ASAP account for the touchdown instances, but not the penalty instances, which is generally what was marked incorrect during our verification process.

A.2 Annotation Event Details

Events for Different Sports In this section, we describe the events that we considered for each sport.

- **Cricket:** Each legal delivery was considered a valid event, where features such as the number of runs and the occurrence of a wide/out ball were marked as well. See Appendix B for further details.
- **American Football:** Each play was considered a valid event, so we considered *punts*, *field goals*, *complete passes*, *incomplete passes*, *run-plays*, *sacks*, *penalties*, and *spikes* as distinct.
- **Football/Soccer:** There are no distinct, sequential plays in football, so we based our events off of online commentary. We mark *shots off target*, *shots on target*, *shots on woodwork*, *goals*, *fouls*, *substitutions*, *yellow cards*, *red cards*, *corner kicks*, *free kicks*, *offsides*, *handballs*, and *saved/blocked balls* as distinct events to be annotated and aligned.
- **Basketball:** Like football/soccer, there are no distinct plays that happen, so we mark *fouls*, *jumper shots*, *layups*, *dunks*, *free throws*, and *regular shots* as distinct events that we annotate and align.

Granularity of Annotations Because the aligned annotations for different sports rely on the timestamps provided by the online commentary source, we observe that different sports are annotated with varying levels of granularity. Thus, when we verify the accuracy of an aligned annotation, we account for these differing levels of granularity with different margins for error. For example, in football, annotations are provided at a minute-level, so if the human annotator marks the event as occurring anywhere outside that range, we consider the annotation to be incorrect; however, for sports like basketball, where annotation timestamps are given by the second, we provide a margin of error of ± 1 second to the timestamp marked by the human. Similar to football, in cricket, an event lasts for 30-40 seconds, so if a human annotator is able to mark the event as occurring anywhere inside that range, we consider the annotation to be correct.

A.3 Raw Videos Source

All of the videos that we ran ASAP through were found on YouTube channels. For cricket we used 131 videos, and for all other three sports we annotated 3 videos each. The average video length of a cricket match is 7.5 hrs while for the other sports it is 1.5 hrs. We also provide the links to all the videos annotated with the supplementary document.

A.4 Qualitative Example

We provide a qualitative sample by attaching an annotation along with a sports (cricket) video snippet. The annotation is present as a `.srt` file and can be used as subtitle with the clip present to see the alignment accuracy of our pipeline.

B LCric

B.1 Primer on Cricket

In this section we further extend our primer to Cricket provided in Section 4.1 by describing the Batting/Bowling phases, as well as the primary objective of the game.

Overview: Cricket is a ball-and-bat sport played by two teams of eleven players each. Cricket is scored by "runs", and at the end of the game, the team with the most scored "runs" wins. The game is played in an inning-format, where one team is batting, and the other team is fielding. We describe the two phases below.

Bowling Phase: When a team is in the bowling phase, all 11 players stay on the field. One of the players is designated as the bowler, and their job is to deliver the ball to the batter (hitter) on the batting team. If the ball is struck by the batsman, the remaining players, called fielders, try to prevent the ball from reaching the boundary of the field and return the ball back to the pitch area. A single over consists of six deliveries bowled by the same player, and each team delivers a set number of overs depending on the tournament type in their bowling phase.

Batting Phase: When is team is in the batting phase, only two players on the team stay on the field at a time. The batsman’s job is to score runs and defend their wickets. A single

run is scored when the batsman hits the ball and runs from one end of the pitch to another. Another way to score runs is to hit the ball to the boundary of the field, which is called the **'boundary'**, giving 4 or 6 runs to the batting team. In total, each batting team has 10 wickets.

Objective: During an inning, the batting team wants to score as many runs as possible, while the bowling team wants to take as many wickets as possible to stop the batting team from scoring. In most single-day matches, the bowling team will bowl for 50 overs before the teams switch roles for the second half of the game. At this point, the goal of the new batting team is to outscore the previous team in runs before 50 overs or before losing all of their wickets.

B.2 Training and Implementation Details

We use consistent training schemes for both TQN (Zhang, Gupta, and Zisserman 2021a) and MeMViT (Wu et al. 2022a) to provide a fair comparison between the two baselines. Both models were trained for 50 epochs on 4 V100 GPUs with a batch size of 4. We used a base learning rate of $LR = 0.01$ with the Adam optimizer and default hyperparameters.

Baseline Implementations For setting up TQN as a baseline, we used the official code provided by the authors with some minor modifications to the output heads for answering **LCric** queries. For MeMViT, since there is no official implementation released at the time of writing, we implemented our own version using the same implementation details as the main paper. Our implementation is built on top of the official implementation of MViT (Fan et al. 2021), which is the base model used to create MeMViT.

B.3 LCric Queries

Query Set Generation Algorithm We describe our query set generation process in Algorithm 1, where we use logical operators and a set of possible atomic events form different combinations of queries.

Binary Queries Statistics For our 10-over experiments, we formed a balanced set of 32 queries by taking queries from the set formed by Algorithm 1 and pruning them down so that given a random 10-over clip sampled uniformly from **LCric**, there would be a 0.5 ± 0.05 probability that the query would hold true on that clip. We list the set of all such queries and their corresponding probabilities in Table ??.

Multi-choice Query Statistics We also generated a set of multi-choice queries for our 10-over experiments. These queries include a mix of common and less common event chains that generally occur between 0-9 (inclusive) times within any 10-over clip. The frequency of occurrence of these clips within our train set is provided in Figure 6.

B.4 AMT Interface

We built an AMT interface for verifying ASAP’s alignment of cricket annotations to videos, with the full instructions and interface provided in Figure 7.

Algorithm 1: Query Set Generation

```

1 # The set of atomic events:  $[0, 1, 2, \dots, 9, W, w]$ 
2 Set of atomic events:  $A_e$ 
3 # The number of queries for the query set
4 Size of the query set:  $n_q$ 
5 query_set = []
6 for  $i$  in  $\text{range}(n_q)$  do
7     # Step A: getting raw operators and combinators
8     # choice
9     num_joins  $\sim [1, 5]$ 
10    # total length for operators set being sampled
11    # for determining the query length
12    ops = random.choices([atleast(), atmost(),
13                        inrange()], num_joins) # sampling list of
14    # operators
15    combine_op = random.choices([and, or], 1) #
16    # sampling the combination operator
17    # Step B: instantiating a query for query set
18    # for pattern matching
19    query = []
20    for op in ops do
21        # specify lower bound for atleast/inrange
22        # ops
23        occ_min  $\sim [1, 10]$ 
24        # specify upper bound bound for
25        # atmost/inrange ops
26        occ_max  $\sim [\text{occ\_min}, 10]$ 
27        # sample atomic events in query
28        atomic_event  $\sim [A_e]$ 
29        # Using the above variables for defining an
30        # occurrence pattern for atomic_event
31        instanced_op = op(occ_min, occ_max, atomic_event)
32        query.append(instanced_op)
33    final_query = join_op(query)
34    query_set.append(final_query)

```

Instruction Details Each annotator is given a set of instructions to read prior to beginning the main annotation task, called a HIT (Human Intelligence Task). For each task, the annotator is given a video clip from a sports match. The task is to classify each legal delivery/ball that occurred in the video, as well as the timestamp at which the annotator was able to gather enough information to answer this question. Additionally, we provide a set of examples for what each event looks like to the annotators, as well as a fully annotated example and video, as shown in Figure 8, 9.

Task Interface Details Each HIT contains a 1-over video and 6 rows, each corresponding to a legal delivery that occurred in the video. Each row consists of a dropdown for inputting the number of runs scored in that delivery, a checkbox for indicating an out ball occurred, a checkbox for indicating a wide ball occurred, and a field for writing the timestamp at which this information can be found. Figure 7 shows what the annotators initially see, as well as an example of how to fill it out.

LCric Annotation Verification A total of 205 overs with 1230 events spanning ~ 1000 minutes were labeled by human annotators and compared to ground truth annotations from ESPNCRicinfo. For each ball, we consider an event annotation to be correct if it was classified completely correctly. The timestamp annotation is marked as correct if it occurred anytime within the timestamp range specified by the ground truth ± 1 seconds.

LCric Annotation Statistics We found that in total, 1185/1230(96.34%) of balls were classified correctly, while 1213/1230(98.62%) of ball timestamps were marked correctly. Additionally, assuming human annotators can aggregate and reason easily with logic, we aggregate their annotations to answer queries in our test set, which provides our human baseline. We find that the human annotations achieve an accuracy of 5541/5740(96.53%) on the test query set – exceeding the TQN and MemViT baselines by a large margin.

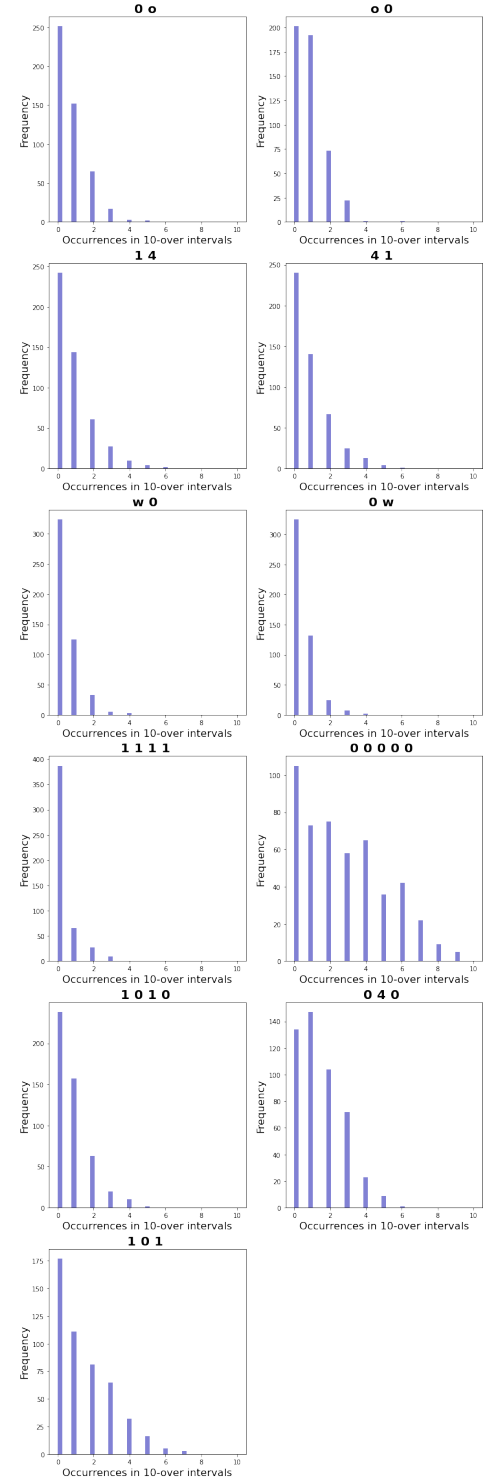


Figure 6: Ground truth output frequencies to queries used in multi-choice queries in the train set of LCric.

Queries	GT probability
atmost 7 1's	0.451
atleast 4 4's	0.523
atleast 5 1's AND atleast 3 4's	0.528
atleast 2 2's AND atleast 3 4's	0.452
atleast 4 4's AND atmost 5 o's	0.452
atleast 4 4's AND atmost 3 5's	0.456
atleast 4 2's OR atmost 2 4's	0.539
atleast 4 3's OR atmost 3 4's	0.544
atleast 5 2's OR atleast 4 4's	0.526
atleast 3 2's OR atleast 2 w's	0.485
atmost 3 4's AND atmost 2 6's	0.529
atmost 3 4's AND atmost 3 7's	0.544
atmost 2 0's OR atmost 3 4's	0.544
2 inrange [1, 6] AND 4 inrange [1, 4]	0.539
4 inrange [1, 6] AND o inrange [1, 4]	0.555
1 inrange [2, 7] OR 2 inrange [4, 5]	0.506
1 inrange [1, 2] OR 2 inrange [2, 3]	0.458
atleast 2 1's AND atleast 2 2's AND atleast 2 4's	0.542
atleast 4 4's OR atleast 4 o's OR atleast 4 w's	0.493
atleast 5 2's OR atleast 4 4's OR atleast 3 6's	0.535
atmost 4 3's AND atmost 3 4's AND atmost 2 5's	0.544
atmost 4 2's AND atleast 3 4's AND atmost 4 w's	0.546
atmost 5 1's OR atleast 5 3's OR atmost 2 4's	0.504
atmost 3 0's OR atleast 5 3's OR atmost 3 4's	0.544
atmost 3 0's OR atmost 4 1's OR atmost 2 4's	0.472
atmost 2 0's OR atmost 5 1's OR atmost 2 4's	0.504
1 inrange [2, 6] OR 2 inrange [3, 4] OR 3 inrange [6, 7]	0.528
atleast 4 0's AND atleast 3 1's AND atleast 2 2's AND atleast 2 4's	0.52
atleast 4 4's OR atleast 2 5's OR atleast 2 6's OR atleast 4 o's	0.518
atmost 3 2's AND atmost 4 4's AND atmost 3 6's AND atmost 5 w's	0.539
6 inrange [1, 7] OR 8 inrange [2, 4] OR o inrange [2, 3] OR w inrange [6, 7]	0.494
1 inrange [1, 6] OR 5 inrange [1, 2] OR o inrange [3, 6] OR w inrange [4, 6]	0.511

Table 4: The binary choice query set used for 10 over experiments and their associated ground truth (GT) probability of occurrence in the **LCric** train set.

Strongly recommended to know the game of Cricket/aware of the rules.

Description

Help us annotate cricket matches by filling in the events happening per ball in a clip.

Instructions

For each cricket match video, there will be up to 6 deliveries that you will need to label. For each delivery, you will need to report:

1. the **number of runs scored** in the delivery
2. whether or not there **was a wide ball or out ball (or neither)** in that delivery
3. **when in seconds** did the batsman play the delivery?

Apart from this, at the very end there is also last question prompt inquiring whether the clip given is sufficient for answering the given set of questions. Please answer it Yes/No accordingly.

Note: If a ball is wide, the ball subsequent to it will also be considered as a part of the same delivery. Also, please do not consider the wide towards the run tally. For instance, if during the second delivery, a bowler bowls a wide ball, then the batter gets 2 runs on the next ball, check "Wide?" and select "2" for the number of runs.

Please find the **timestamp info** for filling out the timestamp related question just above the clip in **red** color.

We request you to watch the full video carefully on a laptop or a computer to precisely answer the questions. The video player has a playback speed option which can be used to alter the playback speed up to 2x.

Please find the detailed instructions below where we cover the process with an example.

We provide an example video with a set of fully labeled annotations. We also walk through how we got each of the annotations labels.

We provide a fully annotated set of labels below for the video above.

Please write down what happened in each legal delivery (up to 6) of 1 over of a cricket match. Select N/A if the legal delivery did not occur.

Time: 285.0



1. Runs?	2	Wide?	<input type="checkbox"/>	Out?	<input type="checkbox"/>	63.7
2. Runs?	1	Wide?	<input type="checkbox"/>	Out?	<input type="checkbox"/>	99.0
3. Runs?	0	Wide?	<input checked="" type="checkbox"/>	Out?	<input type="checkbox"/>	171.7
4. Runs?	0	Wide?	<input type="checkbox"/>	Out?	<input type="checkbox"/>	206.0
5. Runs?	1	Wide?	<input type="checkbox"/>	Out?	<input type="checkbox"/>	247.8
6. Runs?	0	Wide?	<input type="checkbox"/>	Out?	<input type="checkbox"/>	283.0

Within a document, navigate to File > Page setup to switch between pages (the default format) and pageless (the new format). Changes to this setting are document-specific: everyone who interacts with your document will see it, but changing the setting for one document won't impact other documents you own.

For annotating the above match the thinking used is as follows:

1. In the *first* delivery, the batsman hits the ball and begins running, resulting in two runs, **so we mark down 2 in the dropdown "Runs?"**. We note that no out-balls or wide-balls occurred, so we **do not check either box labelled "Wide?" or "Out?"**. We then pause the video at the point when the batsman hit the ball and started running and read the **red timer on the top left of the video** that shows the current time we are paused on, and **mark that time down in seconds [63.7]** in the right-most blank (*do a rough estimate of the time the batsman hit the ball to the best of your ability*).
2. In the *second* delivery, the batsman hits the ball and scores a single run, **so we mark down 1 in the dropdown "Runs?"**. We note that no out-balls or wide-balls occurred, and **write down the time [99.0]** that the batsman hit the ball.
3. "In the *third* delivery, the batsman is first thrown a wide ball. So we check off the **wide-ball** label. Since the batsman was thrown a wide ball, we count the subsequent ball as part of the same delivery. In the next ball, the batsman scores 0 runs, **so we mark down 0 in the dropdown "Runs?"**. We then **mark the time that the batsman hit the ball [171.7]** (you can mark either when the batsman was thrown the wide ball, or when the batsman hit/missed the subsequent ball).
4. In the *fourth* delivery, the batsman misses and scores no runs, **so we mark down 0 in the dropdown "Runs?"**. We then **mark the time that the batsman swung at the ball [206.0]**.
5. In the *fifth* delivery, the batsman hits the ball and scores a single run, **so we mark down 1 in the dropdown "Runs?"**. We note that no out-balls or wide-balls occurred, and **write down the time [247.8]** that the batsman hit the ball.
6. In the *sixth* delivery, the batsman hits the ball and scores no runs, **so we mark down 0 in the dropdown "Runs?"**. We note that no out-balls or wide-balls occurred, and **write down the time [283.0]** that the batsman hit the ball.

Finally, we scroll down and answer the last question. Because we were able to answer all of the given questions using the video, we answer "Yes".

Figure 7: AMT instructions page given to annotators prior to starting the task.

Examples of various different kinds of balls

Below we provide some example snippets of various different kinds of balls that can be seen in the video snippets for our task.

1. Dot ball (where run scored is 0):



As can be seen from the clip, the runs scored in the ball is 0. By definition, this can happen either if the batsman does not hit the ball or if he/she hits the ball but is not able to run from one end of the pitch to another.

2. 1 Run Scored:



As can be seen from the clip, the runs scored in the ball is 1. By definition, if a batsman is able to hit the ball and run from one end of the pitch to another, their team is awarded one run. Similarly, a player can score other possibilities of runs such as 2,3, etc.

3. 4 Run (boundary) Scored:



As can be seen from the clip, the runs scored in the ball is 4. By definition, it happens if the batsman hits the ball and the ball hits the ground before reaching the stadium boundary.

4. 6 Run (boundary) Scored:

Figure 8: Instructions page for AMT interface for Cricket. Each of the 12 events is described in gif format.



As can be seen from the clip, the runs scored in the ball is 6. By definition, it happens if the batsman hits the ball and the ball reaches the stadium boundary without hitting the ground.

5. **Out ball**



As can be seen from the clip, the ball leads to the player getting out. By definition, an out can happen on multiple accounts.

- Leg Before Wicket: If a ball delivery hits any part of the body and is adjusted to have been hitting the stumps.
- Run Out: A batsman is deemed run out if a member of the fielding team puts down the wicket while the batsman is out of their crease/ground.
- Bowled Out: A batsman is considered bowled out if a delivery strikes their wicket and puts it down.
- Caught: If a ball is hit by the batsman is caught by the opposing team before it hits the ground, it is considered an out ball as well.

For this task of annotation, we request you to consider the ball where an Out occurs as one where runs scored is also 0.

6. **Wide ball**



As can be seen from the clip, the ball is a wide one. By definition, a ball is considered wide if it is bowled too wide to be played by a batsman. Also, a wide ball leads to another ball being played on the same ball number and 1 run also being awarded.

Figure 9: Instructions page for AMT interface for Cricket. Each of the 12 events is described in gif format.